

BioJulia


Design and scope

Jakob Nybo Nissen, 2021-10-21

Overview

- What do bioinformaticians do?
- What would the ideal bioinformatics framework look like?
- The values of BioJulia
- BioSequences

Bioinformatics as a discipline



- Biological sequences *are* digital! No, really!
- This makes bioinformatics an "easy" field in some ways compared to most computational biology




Field that develops methods and software tools for understanding biological data



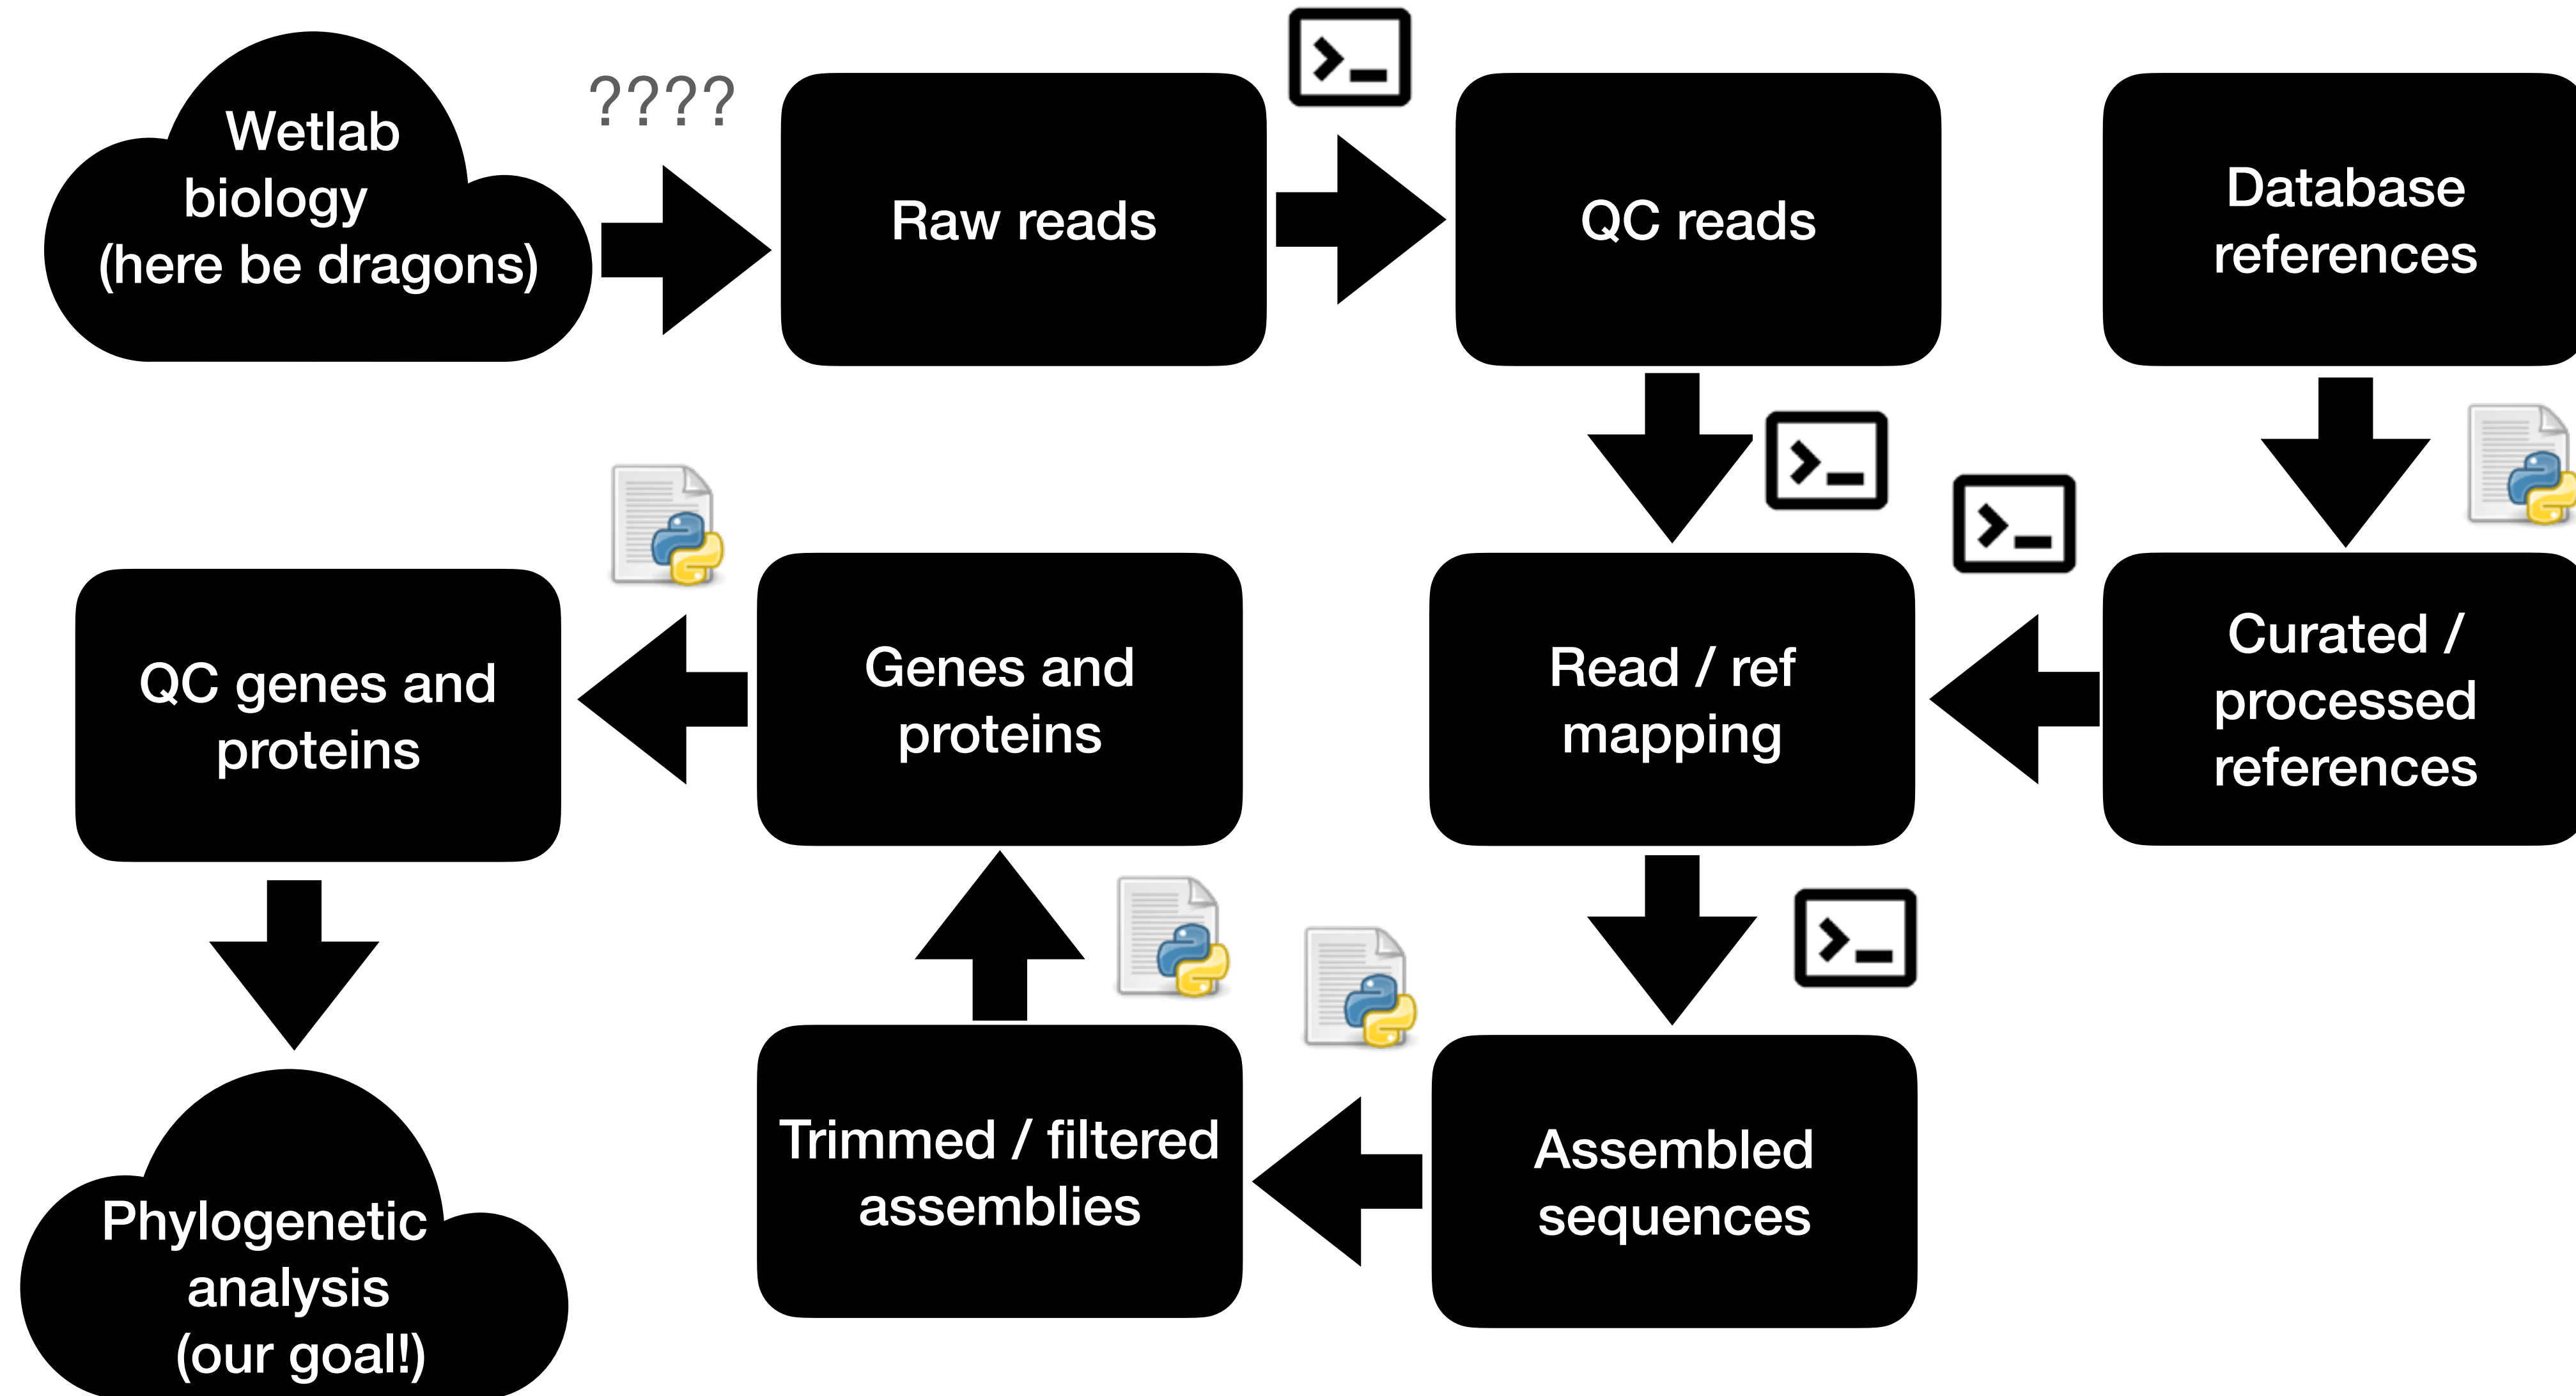
*Field that **use** and develop software for analysing biological **sequences and molecules** as **digital data***

The most abused graph in bioinformatics

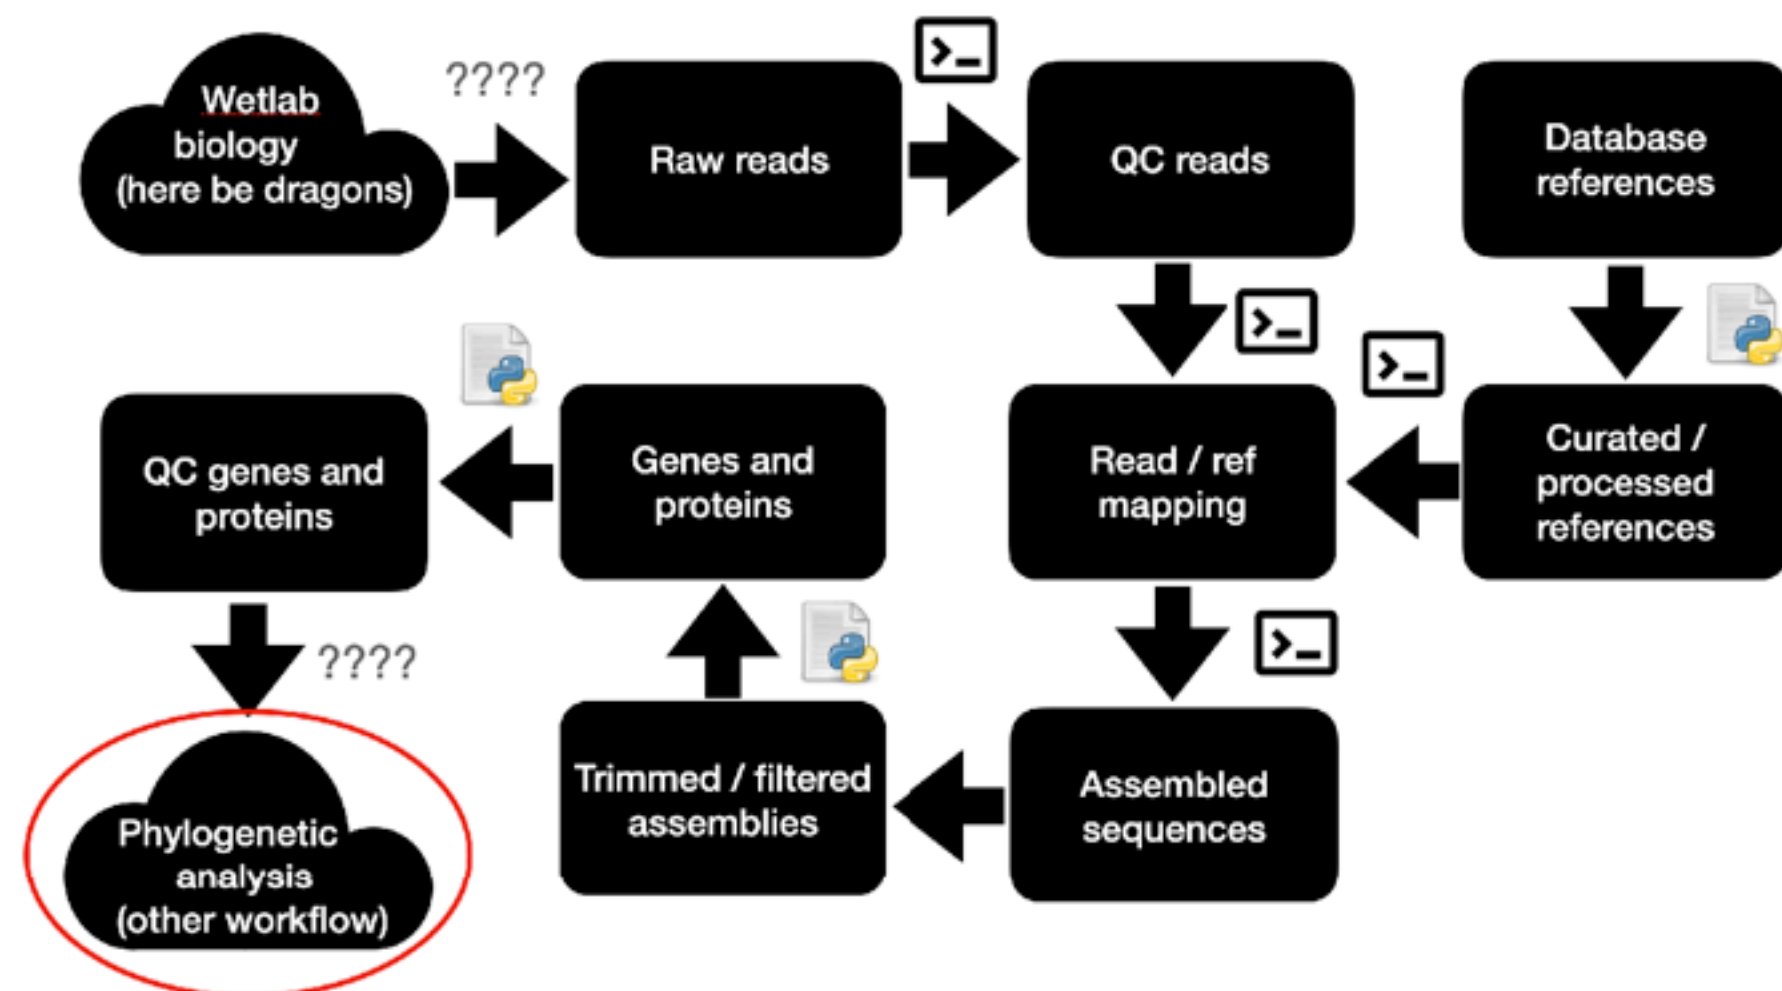


A similar story can be told about mass spectrometry, spatial data, and perhaps microfluidics?

Typical bioinformatics workflow

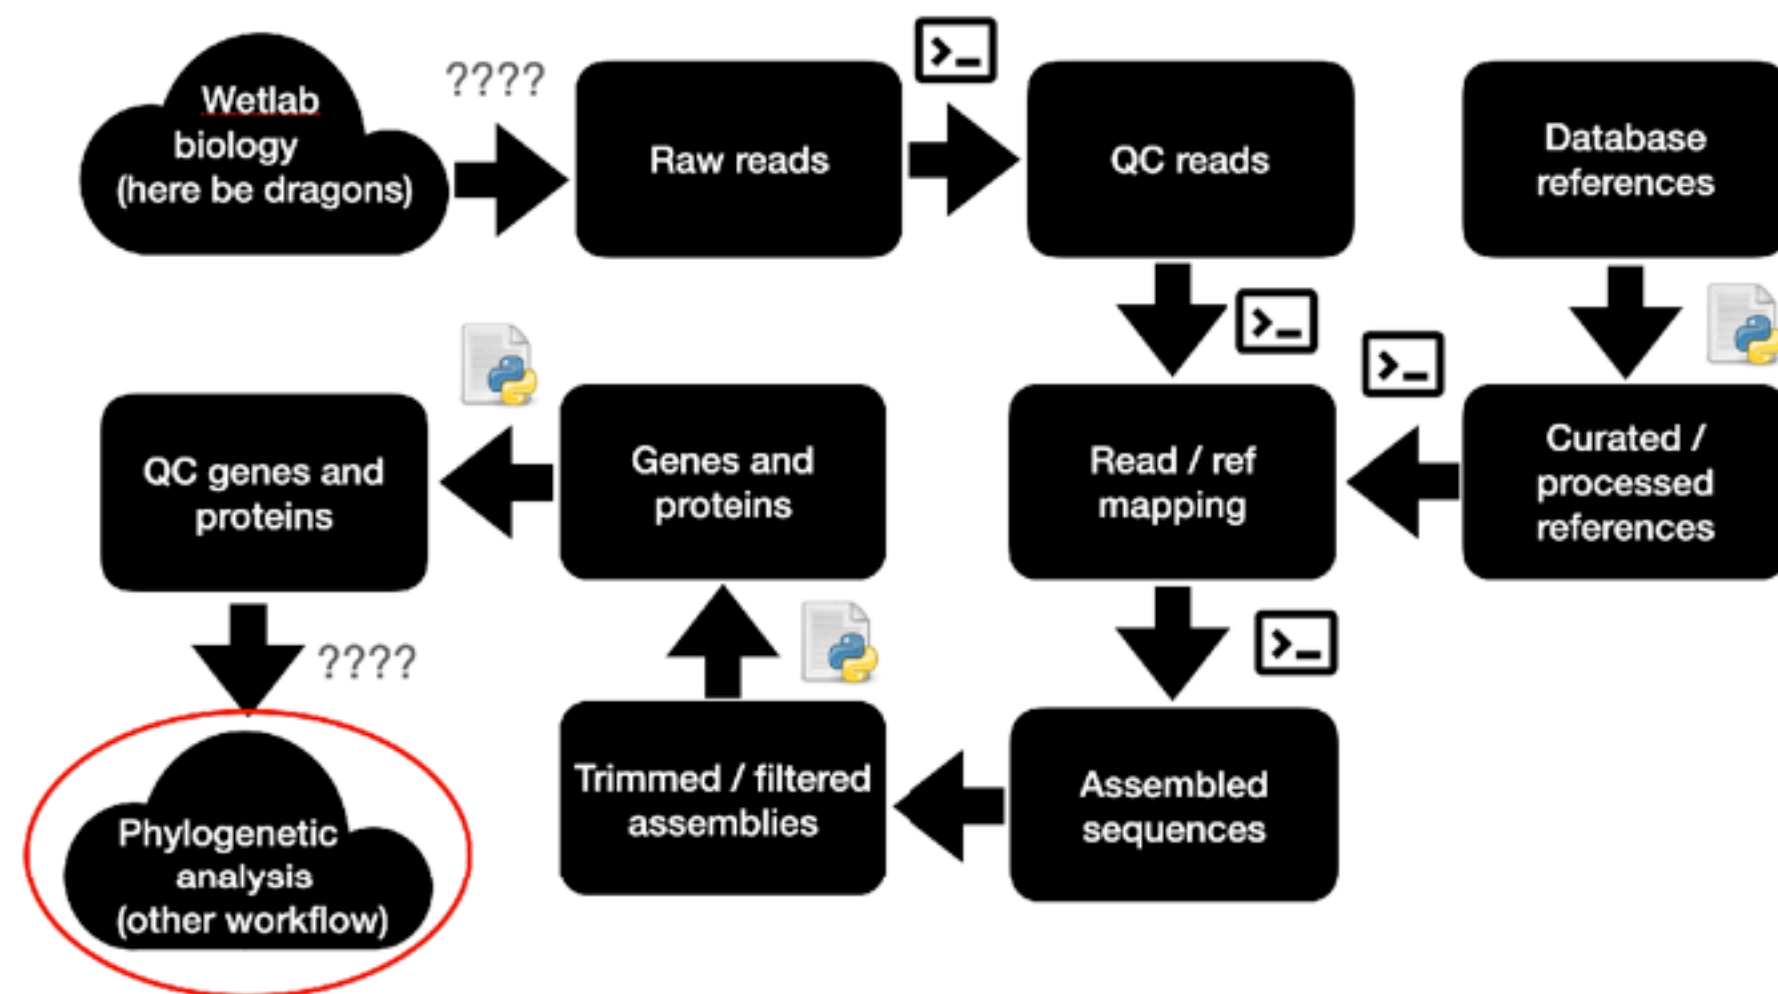


Typical bioinformatics workflow

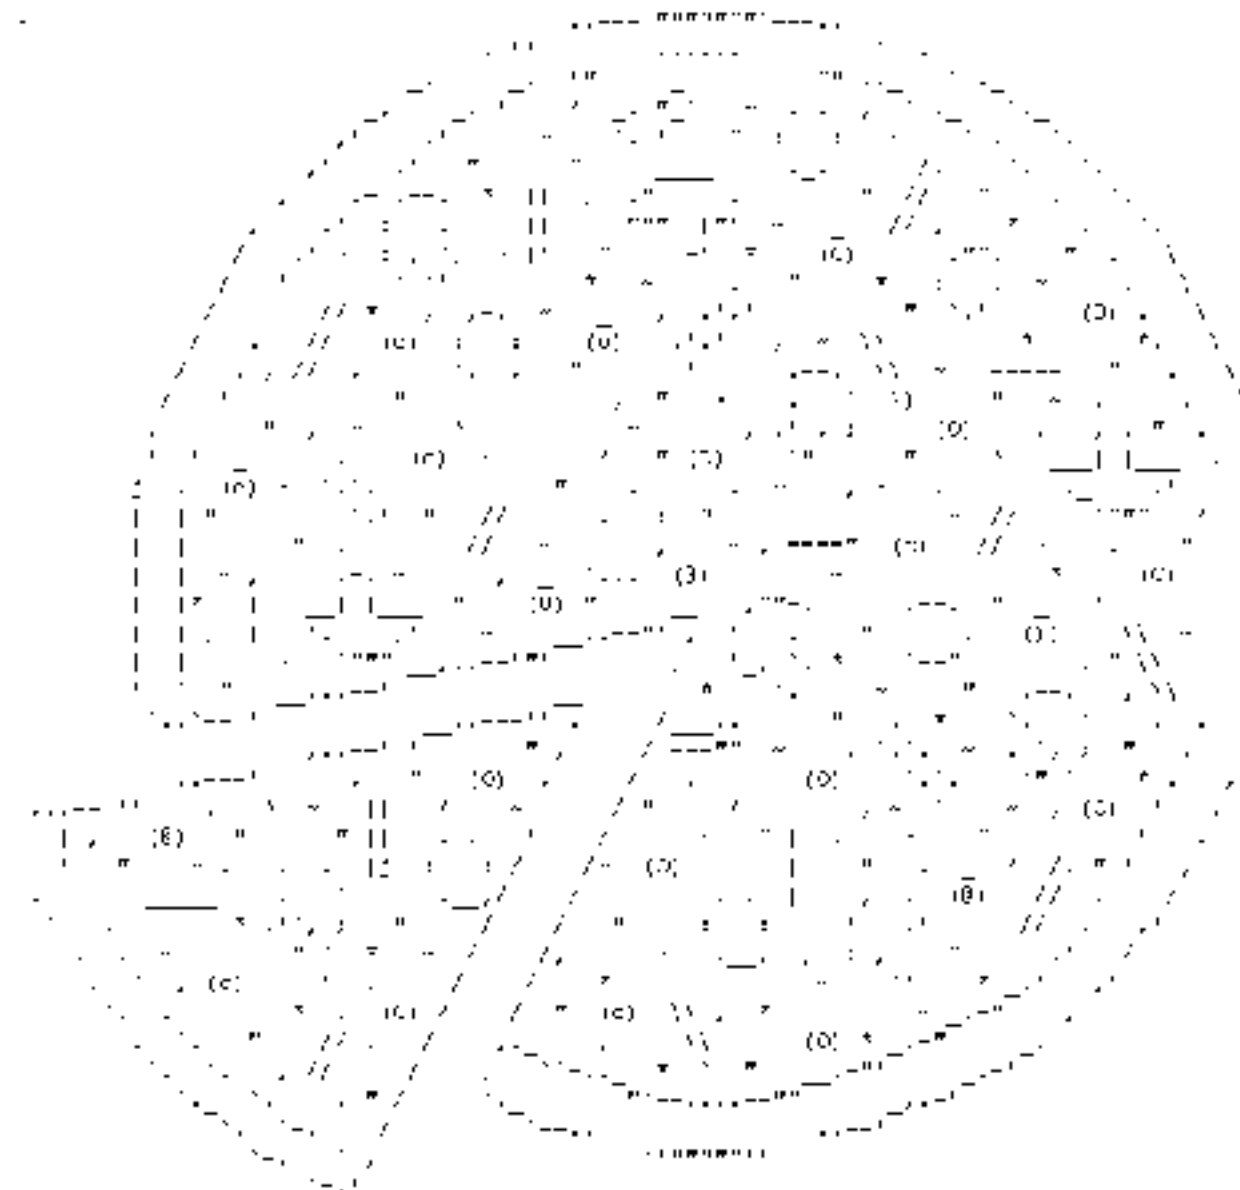


- Workflows are mix of compiled tools and python/R scripts "duct taped" together
- Notoriously unreliable and brittle!
- Many tools *must* be in C/C++ because of the large amounts of data = speed!
- "Bag of tricks" approach all too common!
- But the large amounts of custom code and ever-changing requirements necessitates dynamic scripting languages

Typical bioinformatics workflow



- Each step reads in data from disk, processes, and serialises back
- This leads to a *huge* amount of different file formats, and parsing problems
- Formats are often underspecified and ad-hoc
- Perhaps deliberately?



Valid DNA sequence file, according to commonly used software

What would a solution look like?

- We need a *fast, dynamic* scripting language
- Rewrite ~~it~~ In Julia?
 - Completely unrealistic. Decades of tools, too much inherent knowledge
 - I can't even rewrite my own PhD project, realistically!
- Instead, improve the *glue language* aspect:
 - Fast enough to never need to use another language due to speed
 - Good at calling into shell and interface with other languages
 - Excellent at parsing
- We also need good "workflow management software", but that's another topic....



BioJulia



- Fast enough to not need C for speed
- Dynamic enough to not need Python for scripting
- Implement *basic, commonly used* bioinformatics algorithms
 - These are the "center of the wheel" of many analyses!
- Types and algorithms should be generic enough to be re-usable across a wide range of applications and tools: Don't re-invent the wheel too many times
- Have great file readers/writers
- Be the swiss army knife of bioinformatics
- Incremental replacement of Python/C with Julia.

BioSequences.jl



- Biological sequences *define* bioinformatics. Very important package!



A bioseq is a `str`!"

"ACG"



A bioseq is a `&[u8]!`"



[65, 67, 71]

- A biological sequence is a biological sequence!
- Internal representation is incidental. Abstract it away unless speed is necessary!
- "Parse, don't validate"

BioSequences.jl



BioSequence{Alphabet}



- BioSequence: Abstract type that all concrete biosequence types subtype
- This abstracts over its representation.

- Alphabet: Abstract type all alphabets subtype.
- An alphabet is the set of biological symbols that can be contained in the sequence
- Abstracts over "validation/parsing".
- The Alphabet + BioSequence type decide how data is encoded concretely.



LongSequence{DNAAlphabet{4}}

- LongSequence: Sequence is arbitrary length, mutable, stored on the heap in a Vector. Elements are encoded by bit-packing UInt64.
- You will use another kind of BioSequence next exercise!
- DNAAlphabet{4}: Sequence can contain "-ACMGRSVTWYHKDBN". Each symbol is packed into 4 bits.
- Other pre-defined alphabets include RNAAlphabet{2} ("ACGU") and AminoAcidAlphabet ("ARNDCQEGHILKMFPSTWYVQUBJZX*-")

Things to do with biological sequences

Reverse complementation



Transcription and translation

DNA: TAG CTA GGC TAA ACA TAG

RNA: UGG CUA GGC UAC UCU UAG

AA : W L G Y S STOP

Questions?

Exercise 2