

High-throughput alignments

SAM and BAM

Jakob Nissen, 2021-10-21


Overview

- Previously: Small reads -> assembly -> genome
- But genome assembly is not the most common use of small reads!
 - Genome assembly is computationally very hard
 - It's been studied for decades, still error prone
- Usually, you align reads to an existing genome. Much easier!



What's the problem with using S/W for this?

Remember the most abused slide?



Problems

- S/W algorithm scales with length of reference and query
 - Align to human chromosome 1 - 250M bases!
- Easily millions of reads
- FASTA output format: Completely infeasible!

```
>read
----- [ 250 mill ... ]
>chrom1
NNNNNATTCGGAGTCGTATTAGGGAGAGCGA [ 250 mill ... ]
```

Solution 1: Seed and extend

[illegible]

- Remember how kmers could be represented by a machine integer?
- Matching two kmers is a single CPU instruction, < 1 nanosecond
- Alignment approach: Match kmers between reference and query, and use them as *seeds*.
 - Around each seed, align using S/W
- Most common algorithms are BLAST (1990). BWA (2009?), minimap2 (2018), KMA (2018)
 - All of these use seed & extend

Solution 2: SAM format

- We need a more efficient file format for alignments than FASTA!
- Sequence Alignment/Map format
 - Complicated format - we'll just go through the important parts
- First a header
 - A sequence of header rows
 - Header row: @ + 2 characters + \t + fields
 - Field: 2 characters + : + data

```
@HD      VN:1.6 S0:coordinate
@SQ      SN:SEQHEADER LN:1501
```

Solution 2: SAM format

- Then each alignment on its own line.
- Fields separated by a tab (\t)

- Query identifier
- Sam flags
- Reference identifier
- 1-based leftmost mapping position

• Mapping quality

• CIGAR string

• Ref of next query if query is grouped

• Pos of next query if query is grouped

• Template length




• Query sequence

• Sequence ASCII quality

• Optional fields

But what is this?

We won't go into these

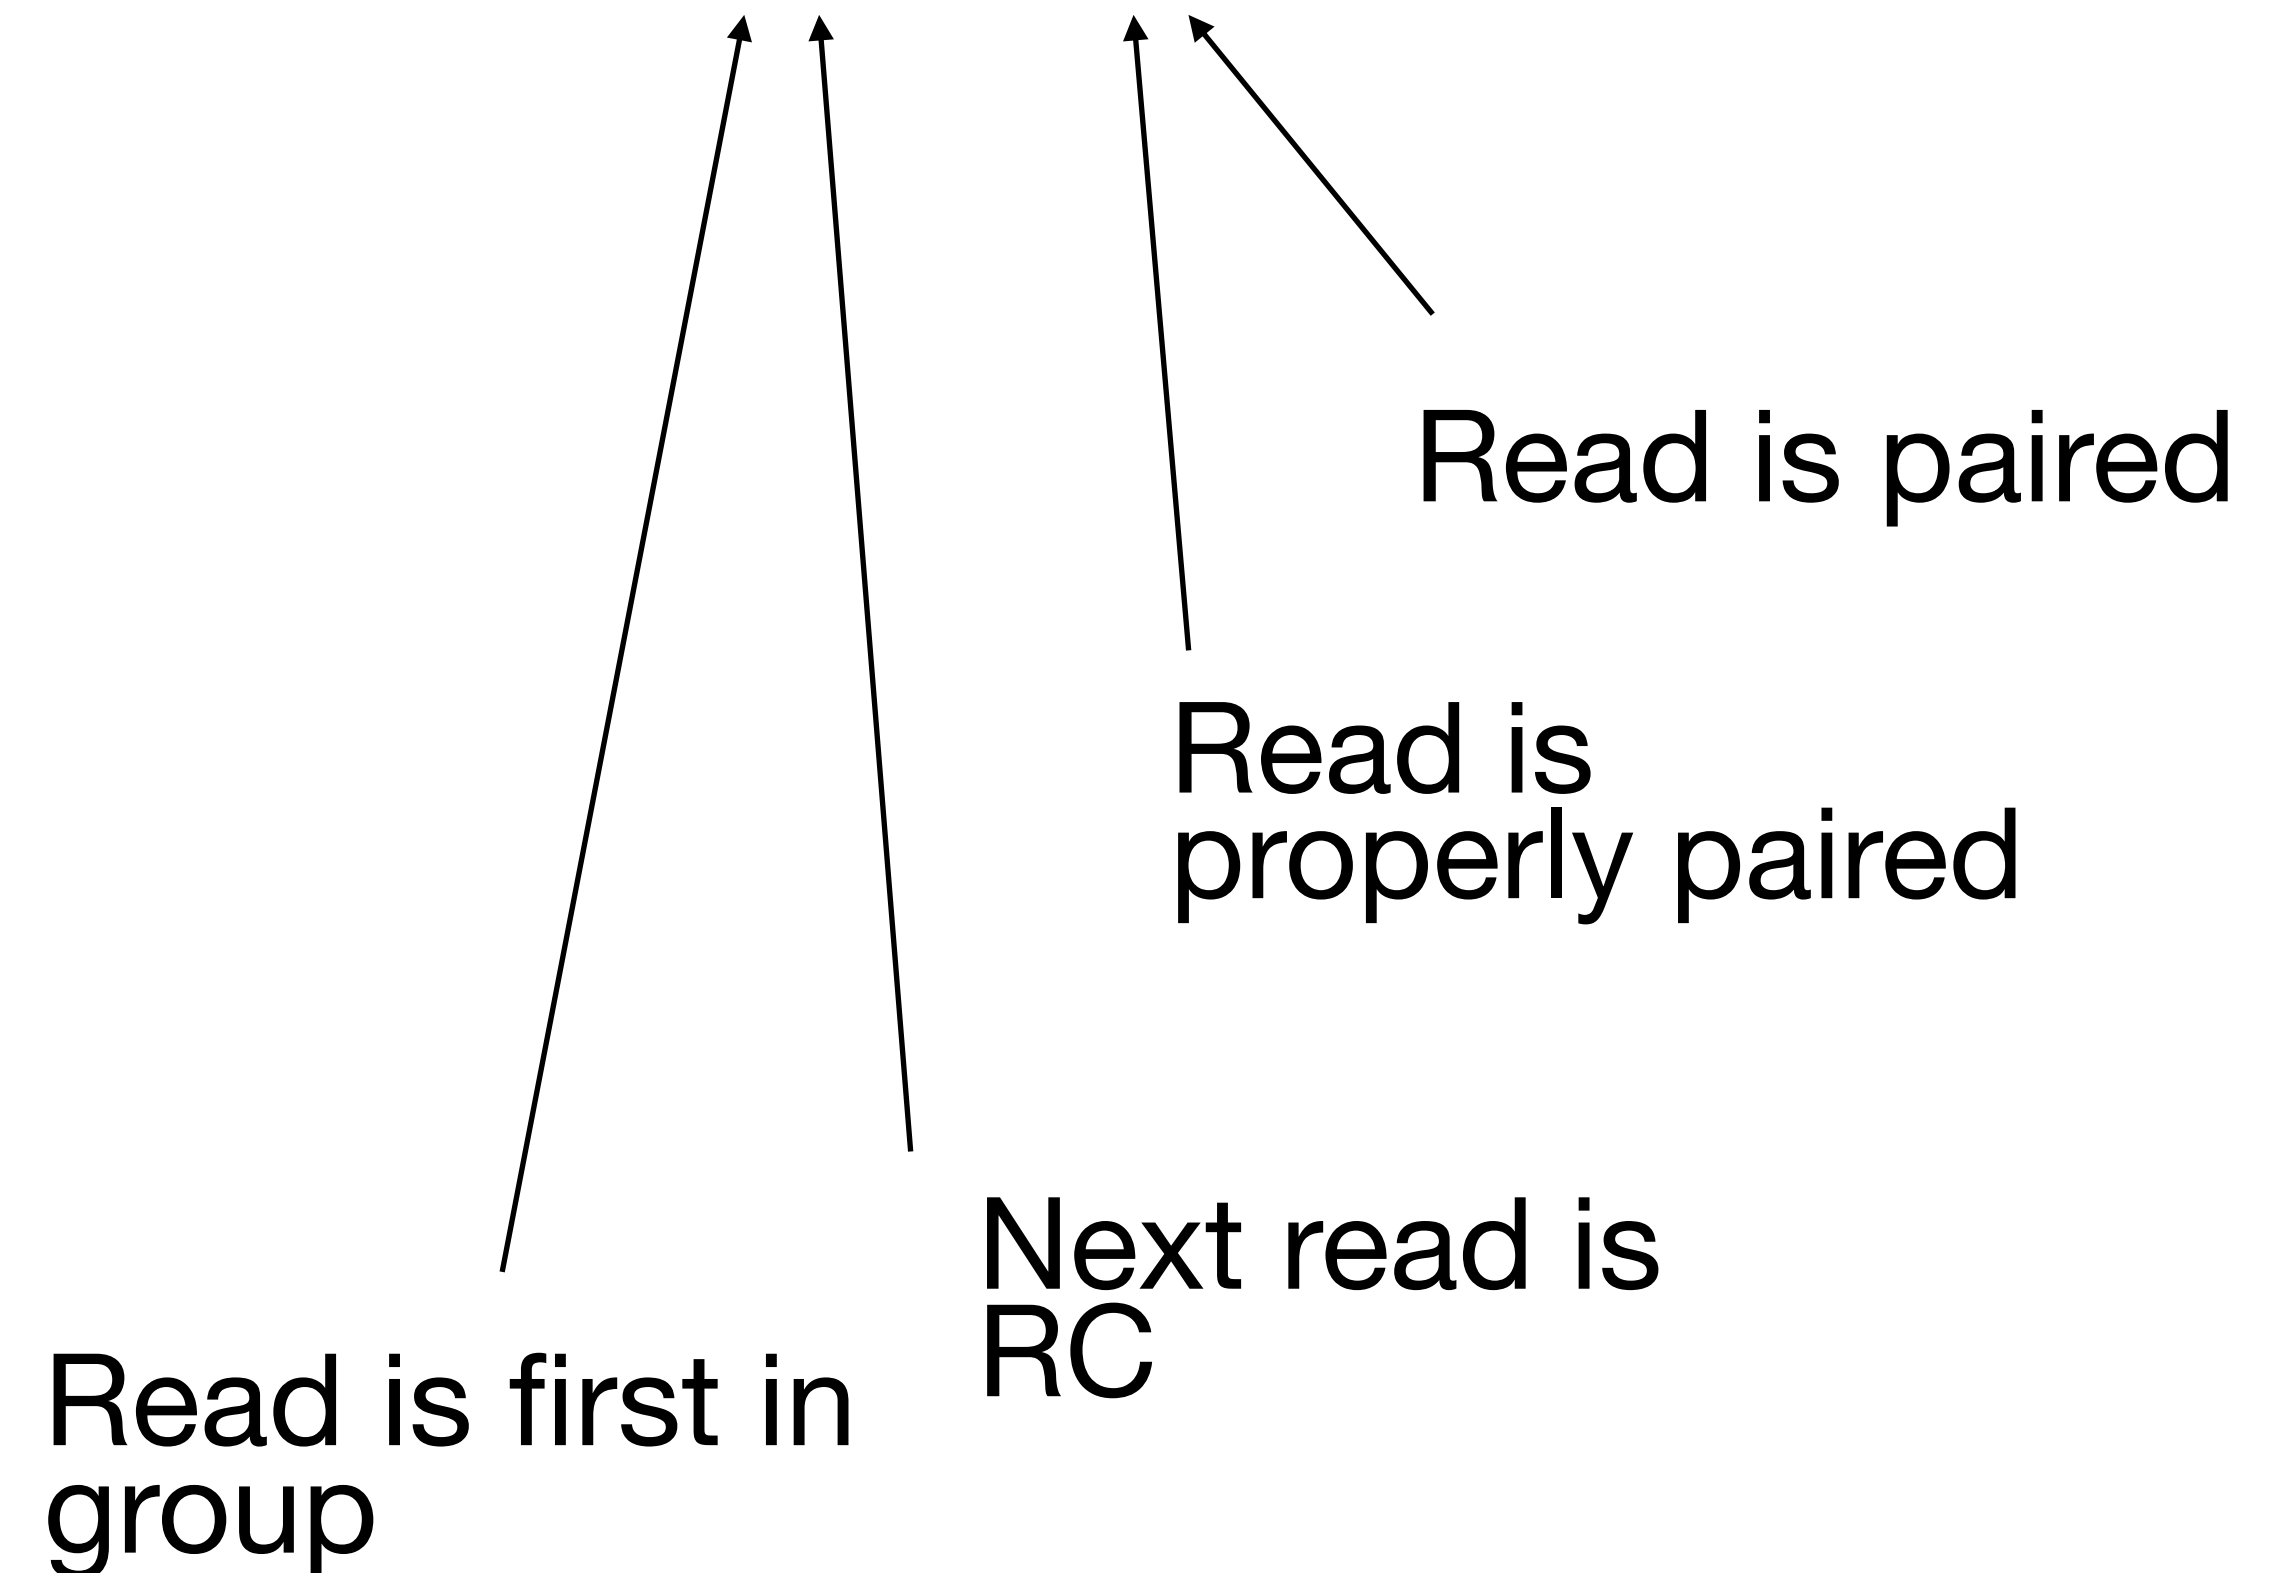


SAM flags

- Each alignment has a list of boolean (true / false) statement
 - Read is paired
 - Read is properly paired
 - Read is not mapped
 - Next read in read group not mapped
 - Read is reverse complemented
 - Next read in read group is reverse complemented
 - Is first read in read group
 - Read is last read in read group
 - Secondary alignment
 - Alignment fails quality checks
 - PCR or optical duplicate
 - Supplementary alignment

Encode it as a number in binary!

0000000001100011 = 99



BAM format

- Binary Alignment Format
- The binary equivalent of SAM - Has a 1:1 correspondence with SAM
- More efficient storage than text
- Compressed in BGZF format
- Everyone uses BAM, no-one uses SAM.
 - Including you, in the exercise

Two language problem...

- BAM files are manipulated with the `samtools` program

```
$ samtools view subset.bam | grep 99 | head -1
```

Languages



● C 72.8%	● Perl 20.3%
● M4 2.5%	● Lua 1.6%
● Shell 1.2%	● Makefile 1.0%
● Other 0.6%	

XAM.jl



- Often you see:
 - Q: "How do I extract all primary alignment that map to this region where..."?
 - A: Use samtools to filter, then pass into grep, then awk, then re-add the header
- This awkwardness doesn't scale to more complicated workflows or really complex read filtering / processing
- XAM.jl: A Julia package for processing / reading / writing SAM and BAM
- XAM.jl is kind of crappy right now, but we will only do basic stuff today
 - We will re-write XAM at some point and make it good

Questions?

Exercise 5